
Processamento de Linguagem Natural

Luís Filipe da Costa Cunha
lfc@di.uminho.pt

José João Almeida
jj@di.uminho.pt

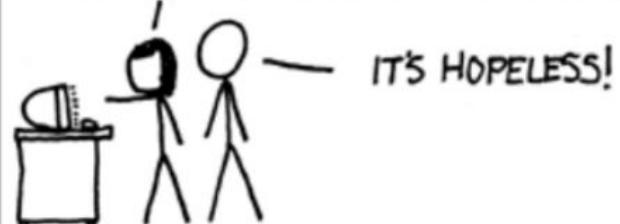


WHENEVER I LEARN A
NEW SKILL I CONCOCT
ELABORATE FANTASY
SCENARIOS WHERE IT
LETS ME SAVE THE DAY.

OH NO! THE KILLER
MUST HAVE FOLLOWED
HER ON VACATION!



BUT TO FIND THEM WE'D HAVE TO SEARCH
THROUGH 200 MB OF EMAILS LOOKING FOR
SOMETHING FORMATTED LIKE AN ADDRESS!



EVERYBODY STAND BACK.



I KNOW REGULAR
EXPRESSIONS.



Expressões Regulares

“Regular expressions are extremely useful in extracting information from text such as code, log files, spreadsheets, or even documents.”

-
- **regular expression ("regex")**: describes a pattern of text
 - can test whether a string matches the expr's pattern
 - can use a regex to search/replace characters in a string
 - very powerful, but tough to read
 - regular expressions occur in many places:
 - text editors (TextPad) allow regexes in search/replace
 - languages: JavaScript; Java Scanner, String split
 - Unix/Linux/Mac shell commands (grep, sed, find, etc.)



In the tranquil woodland, a Woodchuck diligently excavated its burrow, while across the clearing, a pair of woodchucks engaged in a playful chase.

As the sun dipped below the horizon, the Woodchucks gathered together near their burrows, their soft chatter filling the air with a sense of camaraderie.

Each woodchuck embodies the spirit of the wilderness in its own unique way, adding vibrancy to the natural landscape.

Regular Expressions

How can we search for any of there?

- woodchuck
- woodchucks
- Woodchuck
- Woodchucks

regex to the rescue!
[wW]oodchuck





Disjunction and Intervals

[AEIOU]	any uppercase vowel
[12345678]	any digit
alun[oa]	aluno, aluna
[A-Z]	any uppercase letter between A and Z
[a-z]	any lowercase letter between a and z
[0-9]	any digit between 0 and 9
[a-zA-Z0-9]	any letter or digit
[^aeiou]	not a lowercase vowel
[^Ss]	...
[^e^]	...
a^b	...



Expressões Regulares

```
import re

text: "Bruno and Bruna love programming in Python. 2025 will be a great year!!"

regex: 'Bruno'                                regex: '[A-Z] [a-z] [a-z] [a-z] [a-z] '
match:

regex:                                         regex: [^aeiouAEIOU]
match: Bruno ; Bruna                         match:

regex:
match: 2024
```



Character Classes

- \d Digit ([0-9])
- \D not \d
- \w letter digit or underscore ([a-zA-Z0-9_])
- \W not \w
- \s whitespace
- \S not whitespace
- . any character except newline (\n)

```
text: "Bruno loves programming in python. 2025 will be a great year!!"
```

```
pattern:
```

```
match: 2025
```



Anchors

- ^ beginning of line
- \$ end of the line
- \b word boundary
- \B not word boundary

```
text: "Bruno loves programming in python!! 2025 will be a great year to create a program !!"
```

```
regex: ^\d\d\d\d
```

```
match:
```

```
regex:
```

```
match: !! (only at the end)
```

```
regex: program
```

```
match:
```

```
regex:
```

```
match: program (word)
```

Quantifiers

- * 0 or more times
- + 1 or more times
- ? 0 or 1 times.
- {n} exactly n occurrences
- {n, } n or more occurrences
- {n, m} between n and m occurrences



Quantifiers

```
import re
```

Text:

Is this a color or colour?

The class started at February 1, 2020

Javascript is not Java

2023 will be a great year!!!

University of Minho is a great place to learn!

Pattern:

'colou?r'

'[0-9]+'

Java[a-zA-Z]*

[0-9]{2,}

\b[a-z]{1,2}\b

Result:

'color', 'colour'

'1', '2020'

'Javascript', 'Java'

'2023'

'of', 'is' 'a' 'to'

Special characters

\ ^ \$. * + ? () [] { } |

- You can escape them by prefixing a backslash

Solve the woodchuck problem!!

Woodchucks is another name for groundhog!

groundhog | woodchuck



Regex Functions (Python)

```
import re
```

- **match** - Try to apply the pattern at the start of the string
- **search** - Scan through string looking for the first location where the regular expression produces a match
- **findall** - Return a list of all non-overlapping matches in the string
- **sub** - Replace occurrences of the regex pattern
- **split** - Split the source string by the occurrences of the pattern



search / match / findall

```
import re

re.match(r'linha', '02-03-2024, esta linha começa com uma data')
re.match(r'...', '02-03-2024, esta linha começa com uma data')
re.match(r'\d{2}-\d{2}-\d{4}', 'O Carnaval foi no dia 01-03-2024')

re.search(r'\d{2}-\d{2}-\d{4}', 'O Carnaval foi no dia 13-02-2024')
re.search(r'\d{2}-\d{2}-\d{4}', 'O Carnaval foi no dia 13-02-2024 e a Páscoa é dia 17-04-2024')

re.findall(r'\d{2}-\d{2}-\d{4}', 'O Carnaval foi no dia 13-02-2024 e a Páscoa é dia 17-04-2024')
```

Raw String

A raw string considers backslashes as literal characters.

```
text ="Hello,\nI'm a student"
```

```
print(text)
```

Output:

```
Hello,
```

```
I'm a student
```

```
text =r"Hello,\nI'm a student"
```

```
print(text)
```

Output:

```
Hello,\nI'm a student
```

Strings in python can be represented in multiple ways

```
len("\n") #1
```

```
len("\\n") #2
```

```
len(r"\n") #2
```

If you want a Python regular expression object which matches a newline character, then you need a 2-character string, consisting of the backslash character followed by the n character

Capturing groups

- Capturing groups are a way to treat multiple characters as a single unit

```
re.findall(r'(Sra|Sr|Senhora|Senhor)', 'A Senhora Teresa encontrou a Sra. Maria no shopping.' )  
['Senhora', 'Sra']
```

```
re.search(r'alde(ão |ãe |õe)s', 'Os aldeões fizeram uma festa na aldeia' )
```

```
re.search(..., '<span>This is the span content<\span> )
```

- Operators after a capturing group are applied to the whole group pattern: `re.match(r'(go)+', 'gogogogo now!')`



split / sub

```
re.split(r' ','O Carnaval foi no dia 01-03-2022 e a Páscoa é dia 17-04-2022')
['O', 'Carnaval', 'foi', 'no', 'dia', '01-03-2022', 'e', 'a', 'Páscoa', 'é', ...]

re.sub(r'and', '&', 'And Baked Beans and Spam')

re.sub(r'and', '&', 'And Baked Beans and Spam', flags=re.IGNORECASE)
```

Watch Out for The Greediness!

- Use regex to match an HTML tag of the following text:

```
<span> Hello World! <\span>
```

regex: `<.+>`

result: ` Hello World! <\span>`

- Greedy will consume as much as possible
- Making it lazy (non greedy)!

regex: `<.+?>`

result: ``

Errors



Suppose you want to find all the occurrences of the word 'the' in a given text.

Pattern 1: `re.search(r'the', text)`

Error 1: Not matching things that we should have matched (The)

Pattern 2: `re.search(r'[Tt]he', text)`

Error 2: Matching strings that we should not have matched (other, then)

Pattern 3:

`re.search(r'[^a-zA-Z][Tt]he[^a-zA-Z]', text)`

Errors

- In NLP we are always dealing with these kinds of errors.
- Reducing the error rate for an application often involves two antagonistic efforts:
 - **Increasing accuracy or precision** (minimizing false positives)
 - **Increasing coverage or recall** (minimizing false negatives).

Exercises

Define regular expressions to match strings that:

1. have a 't'
2. have a 't' or a 'T'
3. have a letter (and how many)
4. have a digit
5. have a decimal number
6. have a length higher than 3 characters
7. have an 'M' but not an 'm'
8. have a character repeated twice

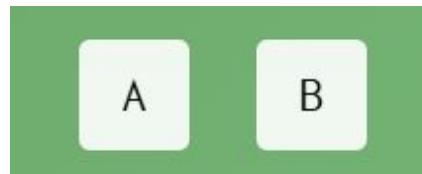
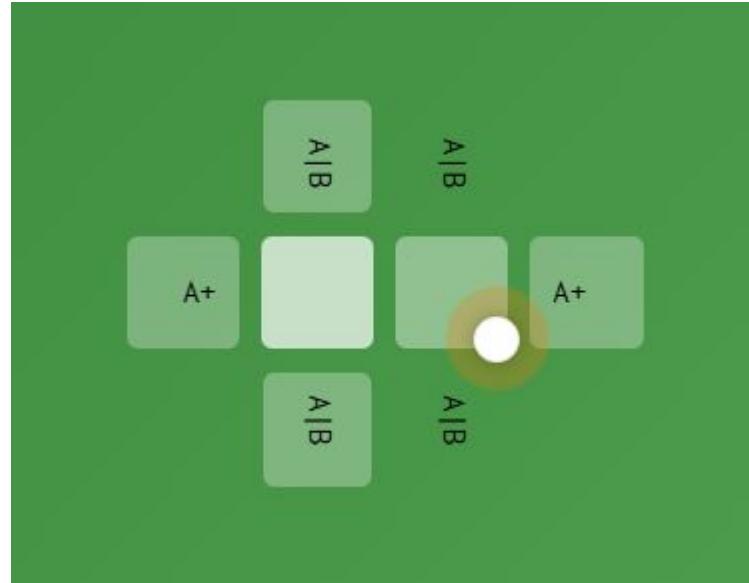


Exercícios

9. Have only one character repeated many times
10. put all words between {}

Regex Cross Word

<https://regexecrossword.com/>





Processamento de Linguagem Natural

Luís Filipe Cunha
[Ifc@di.uminho.pt](mailto:ifc@di.uminho.pt)

José João Almeida
jj@di.uminho.pt

